Uriah Raffer | JC-522616
Statement of Qualifications

1. Describe your experience managing service desk tickets submitted by various customers. How do you ensure tickets are completed on time? How would you identify problems that may be early indicators of a larger issue?

I built and configured a self-hosted GLPI instance in my home lab, an open source IT service management (ITSM) platform in the same category as Cherwell Service Management. I set up an Incident and Service Request workflow with four categories, Hardware, Software, Network, and Access, and I worked each ticket by triaging it, setting its urgency, impact, and priority independently instead of defaulting everything to one severity. A boot failure that fully blocked someone got a high priority, a partial virtual private network (VPN) disruption sat in the middle, and a routine new monitor request stayed low, because the size of the response should match the size of the problem. The people behind tickets sit at very different technical levels, so I adapt how I explain a fix, plain and step by step for a non-technical user and more direct for someone comfortable with the tools. I built that same queue discipline first at Block Advisors, where I carried a caseload of roughly 70 tax returns in a season. Each was a different client working against a hard filing deadline, so managing the queue meant always knowing which ones were closest to their deadline and couldn't be allowed to slip.

I keep tickets on time with the structure I built into GLPI. I set a service level agreement (SLA) with a four hour response target and a two day resolution target on a 24/7 calendar, then added an escalation rule that fires a full day before the resolution deadline on any ticket still open and automatically raises it to the highest priority. The point isn't just setting a deadline. It's surfacing the ticket that's about to miss one while there's still time to act, and I used a priority bump rather than an email alert so the warning shows right in the queue. For a fix that has to be applied the same way more than once, I use a PowerShell script rather than repeating the steps by hand, built and tested in a lab before it goes out, the same way I applied and verified configuration fixes in my internship, so I resolve the same issue the same way every time and don't lose the queue to repeat work.

Catching a larger issue early usually means noticing when one ticket isn't really isolated. When something comes in, I check whether it reproduces for other people, because several users hitting the same symptom points to shared infrastructure rather than a single machine. In a Western Governors University lab a user reported that a normal work site kept loading the wrong page. I tried it from my own machine and had a few others try it. It failed everywhere, which told me the fault was network-wide. I opened Windows Domain Name System (DNS) Manager, found a rogue forward lookup zone redirecting the domain, deleted it, and had the users flush their cache. One report had been the visible edge of a problem hitting everyone. On the systems side, categorizing and routing tickets the same way every time is what makes that pattern surface. In GLPI I set a rule that auto-assigns any Network ticket to the Network team as soon as it's created, and I confirmed it fired instead of assuming the config was right. Once tickets are categorized consistently, a cluster landing in one category in a short window reads as a trend instead of staying buried in the queue, and following up with users on the status of their service tickets is often where that link first shows.